

ASHUTOSH SHARMA

ashutosh14072@gmail.com | +91 9651895627 | [LinkedIn](#) | [GitHub](#)

Generative AI Engineer

PROFESSIONAL SUMMARY

AI Engineer with 3.5+ years building production RAG and agentic systems for enterprise content automation and document intelligence. Delivered 5 end-to-end AI products, reducing operational costs by 50% across 3 major deployments.

KEY HIGHLIGHTS

- **80%+ answer faithfulness** on enterprise RAG through hallucination detection and context filtering, reducing errors by 60%
- **6x ingestion format coverage** by adding JSON, XLSX, CSV, DOCX, and HTML parsers on top of the PDF baseline
- **Sub-second p95 latency** on production query service handling 5,000+ requests daily across distributed infrastructure
- **92% tool-selection accuracy** on agentic search workflows with intelligent routing and error-recovery safeguards

TECHNICAL SKILLS

Generative AI & LLM: LlamaIndex, LangChain, RAG (Hybrid Search, CRAG), Agentic AI, ReAct Agents, Function Calling, Structured Output, Prompt Engineering, Guardrails, RAGAS, NLI Citation Verification, Hugging Face Transformers

Backend & Infrastructure: Python (async), FastAPI, PostgreSQL, pgvector, SQL, Celery, RabbitMQ, WebSocket, REST APIs, Pydantic

Cloud & DevOps: Docker, Kubernetes, Helm, Git, CI/CD, AWS S3, Object Storage, Phoenix Tracing, Prometheus

ML & Data Science: PyTorch, scikit-learn, NLP, Deep Learning, Classification, Clustering, Topic Modeling, Feature Engineering, Pandas

PROFESSIONAL EXPERIENCE

Data Scientist (Generative AI Engineering) | T-Systems ICT India Pvt. Ltd.

December 2022 – Present

AI Operator – Agentic RAG Platform | *May 2025 – Present*

- Design and build a multi-node RAG platform on LlamaIndex and FastAPI, lowering p95 response time by 35% under production load
- Engineer ReAct agent across 6 retrieval tools with loop-prevention safeguards, raising answer faithfulness from 40% to 80%+ via NLI verification and CRAG filtering
- Collaborate with cross-functional teams to extend ingestion from PDF-only to 6 formats (JSON, XLSX, CSV, DOCX, and HTML), onboarding 15+ data sources
- Integrated RAGAS evaluation pipeline with Phoenix tracing for LLM observability; streamlined prompt optimization across 1,000+ weekly evaluations, and presented findings to VP of Engineering

Insurance Healthcare Agent – Agentic RAG System | *December 2024 – April 2025*

- Developed agentic RAG system integrating document and PostgreSQL sources, improving query resolution efficiency by 40%
- Partnered with domain experts to optimize agent workflows using function calling, reducing agent workload by 50% and cutting wait times by 70%
- Spearheaded deployment processing 10,000+ monthly queries at 98% accuracy with resilient fallback handling and 99.5% uptime

Crew AI Multi-Agent Content System | *July 2024 – November 2024*

- Coordinated multi-agent content pipeline using CrewAI with specialized writer, researcher, and editor agents across five distinct use cases
- Accelerated content production by 80% on the blog pipeline while maintaining editorial quality standards

Topic Modeling for Document Clustering | *February 2024 – July 2024*

- Analyzed 1,000+ enterprise documents using LDA topic modeling, boosting search and content organization by 45%
- Trained a LightGBM classifier reaching 90% accuracy on extracted topics, eliminating 70% of manual categorization effort

Enterprise Data Analytics & ML Modeling | *December 2022 – January 2024*

- Directed customer segmentation and classification models using XGBoost and Random Forest, reaching 87% prediction accuracy
- Streamlined ETL pipelines processing 50K+ records with feature engineering, compressing data preparation time by 60%
- Delivered interactive analytics dashboards and reporting pipelines for enterprise clients, improving decision cycles by 30%

Intern Data Scientist | Analytics Labs

November 2021 – November 2022

- Constructed a Random Forest model for customer attrition analysis, driving a 12% reduction in targeted segment attrition
- Established KPIs in the Memory Evaluation Matrix, enabling real-time CPE tracking across 500+ devices

EDUCATION

Master of Science in Data Science | Manipal Academy of Higher Education, Manipal, India | July 2022 – July 2024

Bachelor of Science in Computer Science, Mathematics & Statistics | Bangalore North University, Bangalore, India | July 2018 – November 2021

CERTIFICATIONS

- Generative AI with Large Language Models – DeepLearning.AI / Coursera
- LangChain for LLM Application Development – DeepLearning.AI
- Machine Learning Specialization – Coursera (Andrew Ng)

LANGUAGES

English – Professional Working Proficiency | Hindi – Native